

Lecture 14 — Learning Theory (Part 1 of 2)

Alex Schwing and Matus Telgarsky

March 1, 2018

Theory

CS Theory.

- ▶ Design and analysis of algorithms.
- ▶ Time complexity, space complexity, ...
- ▶ Often worst-case.

Theory

CS Theory.

- ▶ Design and analysis of algorithms.
- ▶ Time complexity, space complexity, ...
- ▶ Often worst-case.

ML Theory.

- ▶ Design and analysis of *ML* algorithms.
- ▶ Time complexity, space complexity, *sample complexity*, *label complexity*, ...
- ▶ Often *average-case*.

Standard ML setup.

- ▶ Want to do well on some *task*; have some input/output pairs.
- ▶ We choose a performance criterion and a family of models.
- ▶ We pick a good model wrt the criterion on the data.

Standard ML setup.

- ▶ Want to do well on some *task*; have some input/output pairs.
- ▶ We choose a performance criterion and a family of models.
- ▶ We pick a good model wrt the criterion on the data.
- ▶ **Examples:** linear regression, logistic regression, SVM, ...
sorting, protein folding, ... !
- ▶ **Counterexamples** (for now): k -nn, ...

Standard ML setup.

- ▶ Want to do well on some *task*; have some input/output pairs.
- ▶ We choose a performance criterion and a family of models.
- ▶ We pick a good model wrt the criterion on the data.
- ▶ **Examples:** linear regression, logistic regression, SVM, ...
sorting, protein folding, ... !
- ▶ **Counterexamples** (for now): k -nn, ...

Formal questions.

- ▶ **Representation/Approximation.** The limitations of our model choice.
- ▶ **Optimization.** Searching for the best model.
- ▶ **Generalization.** Gap between training and testing errors.

Why?

Why?

- ▶ Helps us reason about machine learning.

Why?

- ▶ Helps us reason about machine learning.
- ▶ Helps us improve/make algorithms.

Why?

- ▶ Helps us reason about machine learning.
- ▶ Helps us improve/make algorithms.
- ▶ We are curious.

Why?

- ▶ Helps us reason about machine learning.
- ▶ Helps us improve/make algorithms.
- ▶ We are curious.
- ▶ Math is comfy.

Representation/Approximation.

[Questions so far?]

Representation/Approximation.

[Questions so far?]

Representation failures.

- ▶ Linear functions can fail on even some simple problems.

Representation successes.

- ▶ Polynomial Kernel SVM, RBF Kernel SVM, 2-layer neural nets can fit any continuous function.

Linear functions do not suffice.

Theorem (Minsky-Papert, '69). Consider the 4-point dataset where the corners of the square $\{\pm 1, \pm 1\}$ are labeled with their product. On this data, every linear classifier makes at least 1 error.

Linear functions do not suffice.

Theorem (Minsky-Papert, '69). Consider the 4-point dataset where the corners of the square $\{\pm 1, \pm 1\}$ are labeled with their product. On this data, every linear classifier makes at least 1 error.

- ▶ **Picture “proof”.** [Must pass between two positives and one negative, fail on other negative.]

Linear functions do not suffice.

Theorem (Minsky-Papert, '69). Consider the 4-point dataset where the corners of the square $\{\pm 1, \pm 1\}$ are labeled with their product. On this data, every linear classifier makes at least 1 error.

- ▶ **Picture “proof”.** [Must pass between two positives and one negative, fail on other negative.]
- ▶ **Algebraic proof.**

Consider any halfspace $H := \{x \in R^2 : a^\top x + b \geq 0\}$.

- ▶ If $\{u, -u\} \in H$ for some u , then $b \geq 0$:

$$\begin{aligned} a^\top u + b \geq 0 \quad \wedge \quad a^\top(-u) + b \geq 0 &\implies a^\top(u - u) + 2b \geq 0 \\ &\implies b \geq 0. \end{aligned}$$

- ▶ If some $v \notin H$ and $b \geq 0$, then $-v \in H$:

$$a^\top v + b < 0 \implies a^\top(-v) - b > 0 \implies a^\top(-v) + b > 2b \geq 0.$$

So for any (a, b) , at least one of the two plusses are wrong, or one of the minuses are wrong.

3-layer networks approximate continuous functions.

Theorem. For any continuous $f : [0, 1]^d \rightarrow \mathbb{R}$ and any $\epsilon > 0$, there exists a 3-layer network g with

$$\int_{[0,1]^d} |f(x) - g(x)| \, dx \leq \epsilon.$$

3-layer networks approximate continuous functions.

Theorem. For any continuous $f : [0, 1]^d \rightarrow \mathbb{R}$ and any $\epsilon > 0$, there exists a 3-layer network g with

$$\int_{[0,1]^d} |f(x) - g(x)| \, dx \leq \epsilon.$$

- **Proof** (sketch). First approximate f with a step (piecewise constant) function; then approximate each step function with a 2-layer network (details in lecture).

Optimization

[The second of three analysis topics; questions so far?]

Optimization

[The second of three analysis topics; questions so far?]

- ▶ **The optimization question:**

How do we choose a function which fits the data?

Optimization

[The second of three analysis topics; questions so far?]

- ▶ **The optimization question:**

How do we choose a function which fits the data?

- ▶ In more detail, **(Regularized) Empirical Risk Minimization**

(ERM): We have data $((x_i, y_i))_{i=1}^n$,

predictors \mathcal{F} ,

a regularizer Reg ,

and a performance criterion ℓ ;

We seek to optimize

$$\arg \min_{f \in \mathcal{F}} \widehat{\text{Risk}}(f) + \text{Reg}(f) \quad \text{where} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

Optimization

[The second of three analysis topics; questions so far?]

- ▶ **The optimization question:**

How do we choose a function which fits the data?

- ▶ In more detail, **(Regularized) Empirical Risk Minimization**

(ERM): We have data $((x_i, y_i))_{i=1}^n$,

predictors \mathcal{F} ,

a regularizer Reg ,

and a performance criterion ℓ ;

We seek to optimize

$$\arg \min_{f \in \mathcal{F}} \widehat{\text{Risk}}(f) + \text{Reg}(f) \quad \text{where} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

- ▶ **Example:** Ridge regression; $f_w(x) := w^\top x$ for some w , and

$$\ell(f_w, x, y) = (w^\top x - y)^2 / 2, \quad \text{Reg}(f) := \frac{\lambda}{2} \|w\|^2.$$

We can find w with gradient descent (or “closed form”

$$(X^\top X + \lambda I)^{-1} X^\top y).$$

Ridge regression in more detail.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w} := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Let's consider our three analysis questions.

Also: what is the role of λ ?

Ridge regression in more detail.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w} := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Let's consider our three analysis questions.

Also: what is the role of λ ?

- **Representation:** not just a linear predictor, moreover $\{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$ since

$$\|\hat{w}\|^2 \leq \frac{1}{n\lambda} \|X\hat{w} - y\|^2 + \|\hat{w}\|^2 \leq \frac{1}{n\lambda} \|X0 - y\|^2 + \|0\|^2 \leq \frac{1}{\lambda}.$$

Ridge regression in more detail.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w} := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Let's consider our three analysis questions.

Also: what is the role of λ ?

- **Representation:** not just a linear predictor, moreover $\{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$ since

$$\|\hat{w}\|^2 \leq \frac{1}{n\lambda} \|X\hat{w} - y\|^2 + \|\hat{w}\|^2 \leq \frac{1}{n\lambda} \|X0 - y\|^2 + \|0\|^2 \leq \frac{1}{\lambda}.$$

- **Optimization:** as discussed in class, to achieve accuracy ϵ , gradient descent needs $\frac{\sigma_{\max}(X) + \lambda}{\sigma_{\min}(X) + \lambda} \ln(1/\epsilon)$ iterations.

Ridge regression in more detail.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w} := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Let's consider our three analysis questions.

Also: what is the role of λ ?

- **Representation:** not just a linear predictor, moreover $\{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$ since

$$\|\hat{w}\|^2 \leq \frac{1}{n\lambda} \|X\hat{w} - y\|^2 + \|\hat{w}\|^2 \leq \frac{1}{n\lambda} \|X0 - y\|^2 + \|0\|^2 \leq \frac{1}{\lambda}.$$

- **Optimization:** as discussed in class, to achieve accuracy ϵ , gradient descent needs $\frac{\sigma_{\max}(X) + \lambda}{\sigma_{\min}(X) + \lambda} \ln(1/\epsilon)$ iterations.
- **Generalization:** *coming up next!*

Overfitting and generalization.

[Final analysis topic; questions so far?]

Overfitting and generalization.

[Final analysis topic; questions so far?]

We train on some data because that is what we have.

But what we *want* is good *future* performance.

- ▶ **Generalizing** means similar past and future performance.
- ▶ **Overfitting** means (vastly) better past performance.

Overfitting and generalization.

[Final analysis topic; questions so far?]

We train on some data because that is what we have.

But what we *want* is good *future* performance.

- ▶ **Generalizing** means similar past and future performance.
- ▶ **Overfitting** means (vastly) better past performance.

Example: Suppose data $((x_i, y_i))_{i=1}^n$, with x_i random, and

$$y_i := \bar{w}^\top x_i + \xi_i$$

with independent and zero mean $(\xi_i)_{i=1}^n$.

- ▶ Algo 1: ordinary least squares.
- ▶ Algo 2: fit an n -degree polynomial.

Overfitting and generalization.

[Final analysis topic; questions so far?]

We train on some data because that is what we have.

But what we *want* is good *future* performance.

- ▶ **Generalizing** means similar past and future performance.
- ▶ **Overfitting** means (vastly) better past performance.

Example: Suppose data $((x_i, y_i))_{i=1}^n$, with x_i random, and

$$y_i := \bar{w}^\top x_i + \xi_i$$

with independent and zero mean $(\xi_i)_{i=1}^n$.

- ▶ Algo 1: ordinary least squares.
- ▶ Algo 2: fit an n -degree polynomial.

Second should *overfit*.

To formalize this, need a *model* for unseen data.

Models for data.

Learning theory provides *many* candidate models.

- ▶ **Statistical** setting: training data and future data drawn IID (independent and identically distributed) from some distribution. (An *average case* setting.)
- ▶ **Online** setting: an adversary constructs examples with full knowledge of what we are doing. (A *worst case* setting.)

Which one is more realistic?

Models for data.

Learning theory provides *many* candidate models.

- ▶ **Statistical** setting: training data and future data drawn IID (independent and identically distributed) from some distribution. (An *average case* setting.)
- ▶ **Online** setting: an adversary constructs examples with full knowledge of what we are doing. (A *worst case* setting.)

Which one is more realistic?

Example: spam filtering.

- ▶ On the one hand, spammers observe what google does, try to break its detection.
- ▶ On the other hand, they don't know exactly what google is doing.

Models for data.

Learning theory provides *many* candidate models.

- ▶ **Statistical** setting: training data and future data drawn IID (independent and identically distributed) from some distribution. (An *average case* setting.)
- ▶ **Online** setting: an adversary constructs examples with full knowledge of what we are doing. (A *worst case* setting.)

Which one is more realistic?

Example: spam filtering.

- ▶ On the one hand, spammers observe what google does, try to break its detection.
- ▶ On the other hand, they don't know exactly what google is doing.

If each morning google re-trains on the last 30 days of data, assuming it is IID, is it so bad? (Mixture of both settings.)

The *Statistical Learning Theory* setting.

- ▶ We receive n examples IID from some underlying distribution.
- ▶ We would like to do well according to some performance criterion *in expectation*.

Example. Least squares: we receive $((x_i, y_i))_{i=1}^n$, we select \hat{w} , and the quantity we want to minimize is

$$\mathbb{E}(\hat{w}^\top x - y)^2,$$

but we can't compute this!

(We only have a finite sample.)

Empirical Risk Minimization and generalization.

We have advocated

(Regularized) Empirical Risk Minimization (ERM):

We choose a function class \mathcal{F} , a loss function ℓ , a regularization scheme Reg , and approximately optimize

$$\operatorname{argmin}_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \text{Reg}(f).$$

Why should this work?

Empirical Risk Minimization and generalization.

We have advocated

(Regularized) Empirical Risk Minimization (ERM):

We choose a function class \mathcal{F} , a loss function ℓ , a regularization scheme Reg , and approximately optimize

$$\operatorname{argmin}_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \text{Reg}(f).$$

Why should this work?

- ▶ Vague intuition: law of large numbers (LLN).
- ▶ Rigorous version: suppose \hat{f} selected *without* using $((x_i, y_i))_{i=1}^n$. If $((x_i, y_i))_{i=1}^n$ are IID, then so are $(\ell(\hat{f}, x_i, y_i))_{i=1}^n$, and by LLN

$$\frac{1}{n} \sum_{i=1}^n \ell(\hat{f}, x_i, y_i) \rightarrow \mathbb{E}(\ell(\hat{f}, X, Y)) \quad \text{as } n \rightarrow \infty.$$

Empirical Risk Minimization and generalization.

We have advocated

(Regularized) Empirical Risk Minimization (ERM):

We choose a function class \mathcal{F} , a loss function ℓ , a regularization scheme Reg , and approximately optimize

$$\operatorname{argmin}_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \text{Reg}(f).$$

Why should this work?

- ▶ Vague intuition: law of large numbers (LLN).
- ▶ Rigorous version: suppose \hat{f} selected *without* using $((x_i, y_i))_{i=1}^n$. If $((x_i, y_i))_{i=1}^n$ are IID, then so are $(\ell(\hat{f}, x_i, y_i))_{i=1}^n$, and by LLN

$$\frac{1}{n} \sum_{i=1}^n \ell(\hat{f}, x_i, y_i) \rightarrow \mathbb{E}(\ell(\hat{f}, X, Y)) \quad \text{as } n \rightarrow \infty.$$

Two issues.

- ▶ We want a bound for finite n .
- ▶ We used $((x_i, y_i))_{i=1}^n$ to select \hat{f} !

Issue #1: bounds for finite samples.

The easiest tool here is **Hoeffding's inequality**.

Theorem (Hoeffding's inequality). Suppose each draw from the distribution lies in the interval $[a, b]$. With probability at least $1 - \delta$ over an iid draw of $(z_i)_{i=1}^n$,

$$\mathbb{E}Z \leq \frac{1}{n} \sum_{i=1}^n z_i + (b - a) \sqrt{\frac{\ln(1/\delta)}{2n}}.$$

- **Remark** (on terminology). This is sometimes called a *concentration inequality*, or a *deviation bound*.

Issue #1: bounds for finite samples.

The easiest tool here is **Hoeffding's inequality**.

Theorem (Hoeffding's inequality). Suppose each draw from the distribution lies in the interval $[a, b]$. With probability at least $1 - \delta$ over an iid draw of $(z_i)_{i=1}^n$,

$$\mathbb{E}Z \leq \frac{1}{n} \sum_{i=1}^n z_i + (b - a) \sqrt{\frac{\ln(1/\delta)}{2n}}.$$

- ▶ **Remark** (on terminology). This is sometimes called a *concentration inequality*, or a *deviation bound*.
- ▶ **Example**. Consider fixed \hat{f} , and binary loss $\ell(f, x, y) := \mathbf{1}[f(x) \neq y] \in [0, 1]$. Then, with probability at least $(1 - \delta)$ over an IID draw $((x_i, y_i))_{i=1}^n$,

$$\Pr [\hat{f}(X) \neq Y] = \mathbb{E} \mathbf{1} [\hat{f}(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbf{1} [\hat{f}(x_i) \neq y_i] + \sqrt{\frac{\ln(1/\delta)}{2n}}$$

- ▶ **Remark** (scaling): Randomly receive n $\$d$ bills; confidence interval scales with $d \sqrt{\ln(1/\delta)/2n}$.

Issue #2: *uniform* deviations.

Previous bounds were for a *fixed function* \hat{f} .

Let's take a step back and use our intuition.

Issue #2: *uniform* deviations.

Previous bounds were for a *fixed function* \hat{f} .

Let's take a step back and use our intuition.

Intuition: larger function classes overfit more.

Issue #2: *uniform* deviations.

Previous bounds were for a *fixed function* \hat{f} .

Let's take a step back and use our intuition.

Intuition: larger function classes overfit more.

Rigorous bound: Suppose we select f from functions \mathcal{F} using data $((x_i, y_i))_{i=1}^n$. Then (with probability at least $1 - \delta$),

$$\mathbb{E} \ell(f, X, Y) \leq \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

Issue #2: *uniform* deviations.

Previous bounds were for a *fixed function* \hat{f} .

Let's take a step back and use our intuition.

Intuition: larger function classes overfit more.

Rigorous bound: Suppose we select f from functions \mathcal{F} using data $((x_i, y_i))_{i=1}^n$. Then (with probability at least $1 - \delta$),

$$\mathbb{E}\ell(f, X, Y) \leq \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

Example.

- ▶ Linear classification $w \mapsto \text{sgn}(w^\top x)$ with $w \in R^d$ has

$$\begin{aligned} \Pr \left[\text{sgn}(w^\top x) \neq y \right] &\leq \frac{1}{n} \sum_{i=1}^n \mathbb{1} \left[\text{sgn}(w^\top x_i) \neq y_i \right] \\ &\quad + \tilde{O} \left(\sqrt{(d + \ln(1/\delta))n} \right). \end{aligned}$$

Issue #2: *uniform* deviations.

Previous bounds were for a *fixed function* \hat{f} .

Let's take a step back and use our intuition.

Intuition: larger function classes overfit more.

Rigorous bound: Suppose we select f from functions \mathcal{F} using data $((x_i, y_i))_{i=1}^n$. Then (with probability at least $1 - \delta$),

$$\mathbb{E}\ell(f, X, Y) \leq \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

Example.

- ▶ Linear classification $w \mapsto \text{sgn}(w^\top x)$ with $w \in R^d$ has

$$\begin{aligned} \Pr \left[\text{sgn}(w^\top x) \neq y \right] &\leq \frac{1}{n} \sum_{i=1}^n \mathbb{1} \left[\text{sgn}(w^\top x_i) \neq y_i \right] \\ &\quad + \tilde{O} \left(\sqrt{(d + \ln(1/\delta))n} \right). \end{aligned}$$

- ▶ **Remark:** How to minimize rhs?

Issue #2: *uniform* deviations.

Previous bounds were for a *fixed function* \hat{f} .

Let's take a step back and use our intuition.

Intuition: larger function classes overfit more.

Rigorous bound: Suppose we select f from functions \mathcal{F} using data $((x_i, y_i))_{i=1}^n$. Then (with probability at least $1 - \delta$),

$$\mathbb{E}\ell(f, X, Y) \leq \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

Example.

- ▶ Linear classification $w \mapsto \text{sgn}(w^\top x)$ with $w \in R^d$ has

$$\begin{aligned} \Pr \left[\text{sgn}(w^\top x) \neq y \right] &\leq \frac{1}{n} \sum_{i=1}^n \mathbb{1} \left[\text{sgn}(w^\top x_i) \neq y_i \right] \\ &\quad + \tilde{O} \left(\sqrt{(d + \ln(1/\delta))n} \right). \end{aligned}$$

- ▶ **Remark:** How to minimize rhs?
Logistic upper bounds...

Simple example: finite classes

When $|\mathcal{F}| < \infty$, can do $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$.

Theorem. With probability at least $1 - \delta$,

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln(|\mathcal{F}|) + \ln(1/\delta)}{2n}}.$$

Simple example: finite classes

When $|\mathcal{F}| < \infty$, can do $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$.

Theorem. With probability at least $1 - \delta$,

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln(|\mathcal{F}|) + \ln(1/\delta)}{2n}}.$$

► **Proof.** Define $\epsilon := \sqrt{\ln|\mathcal{F}|/\delta/2n}$ and events

$$E_j := \left[\Pr[f_j(X) \neq Y] > \epsilon + \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f_j(x_i) \neq y_i] \right].$$

By Hoeffding, $\Pr[E_j] \leq \delta/|\mathcal{F}|$, and by union bound

$$\Pr[\cup_j E_j] \leq \sum_j \Pr[E_j] \leq \delta.$$

Where to go from here?

Have much more sophisticated bounds of the form

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right),$$

where:

- ▶ $\text{Risk}(f) = \Pr[f(X) \neq Y]$, \mathcal{F} is ReLU networks with p parameters and L layers, $\text{Complexity}(\mathcal{F}) = \tilde{O}(pL)$.
- ▶ $\text{Risk}(f)$ is least squares risk, $\mathcal{F} := \{w \in \mathbb{R}^d : \|w\| \leq \sqrt{1/\lambda}\}$ (as in Ridge regression), $\text{Complexity}(\mathcal{F}) = 1/\lambda$.

Summary

- ▶ **ML Theory** is design and analysis of ML algorithms.

Summary

- ▶ **ML Theory** is design and analysis of ML algorithms.
- ▶ **Standard setup:** we fit a model to data according to some performance criterion.
 - ▶ **Representation:** is our model powerful enough to capture this phenomenon?
 - ▶ **Optimization:** can we (efficiently?) fit our model to data?
 - ▶ **Generalization:** does our model perform well on unseen data?

Summary

- ▶ **ML Theory** is design and analysis of ML algorithms.
- ▶ **Standard setup:** we fit a model to data according to some performance criterion.
 - ▶ **Representation:** is our model powerful enough to capture this phenomenon?
 - ▶ **Optimization:** can we (efficiently?) fit our model to data?
 - ▶ **Generalization:** does our model perform well on unseen data?
- ▶ **Generalization/overfitting** requires a model of *unseen* data.
 - ▶ **Today** we sketched the *statistical learning theory setting*.
 - ▶ **Next time** we'll go into it in more detail.