

Lecture 15 — Learning Theory (Part 2 of 2)

Alex Schwing and Matus Telgarsky

March 6, 2018

Schedule for today.

- ▶ Midterm announcements.
- ▶ Overfitting/generalization: reminders.
- ▶ Overfitting/generalization: intuition and basics.
- ▶ Overfitting/generalization: some asides.

Learning theory reading:

see my learning theory course's resources (click on <http://mjt.cs.illinois.edu/courses/mlt-f17/> or google "matus uiuc mlt-f17").

Midterm announcements.

- ▶ **Location (all ECEB):** your netid determines your room:
 - ▶ **1002** (this room): aa18 - ryang28.
 - ▶ **1013:** sabag2 - xunlin2.
 - ▶ **1015:** xyu69 - zzhou51.
- ▶ **Time:** start time 6pm, duration 90 minutes.
- ▶ **Notes:** can bring one standard size (“US letter”) sheet of notes, front and back, **handwritten**.
- ▶ **Review lecture and materials:** wait until Thursday.

Learning Theory (part 2 of 2) – Overfitting/Generalization.

- ▶ **Overfitting:** better performance on past data than future data.
- ▶ **Generalizing:** similar performance on past and future data.

Learning Theory (part 2 of 2) – Overfitting/Generalization.

- ▶ **Overfitting:** better performance on past data than future data.
- ▶ **Generalizing:** similar performance on past and future data.

Models

Reasoning about this requires a **model** linking past and future.

- ▶ **Statistical learning theory:** past and future examples drawn IID from a common distribution.
- ▶ **Online learning:** adversary constructs new examples.

Overfitting by example.

Linear or polynomial least squares?


Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.

$\hat{R}_1 = 0.00692304, R_1 = 0.00897457, \hat{R}_{21} = 0.0026458, R_{21} = 1.24795e + 07$




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.

$\hat{R}_1 = 0.00692304, R_1 = 0.00897457, \hat{R}_{34} = 0.00259486, R_{34} = 3.97751e + 10$




Overfitting by example.

Linear or polynomial least squares?

Truth: $y = 0 \cdot x + \xi$, $\xi \sim \text{Gaussian}$.

Red: seen data. **Green:** unseen data.

$\hat{R}_1 = 0.00692304, R_1 = 0.00897457, \hat{R}_{38} = 0.00242094, R_{38} = 1.02285e + 13$



Overfitting by example.




Figure 1: Fitting a degree 38 polynomial.

Intuition: More parameters \implies more overfitting.

Can reduce overfitting with model choice and regularization (similar...).

What learning theory gives us: concrete relationships.

Aside/review: least squares code for this plot!

```
# [ . . . ]

for s in [ 'tr', 'te', 'grid' ]:
    X[s][:, 0] = 1.0
    #X[s][:, 1] is random according to some distribution
    for j in range(2, n):
        X[s][:, j] = X[s][:, 1] * X[s][:, j - 1]

# [ . . . ]

for j in range(2, n):
    #better to use the black box N.linalg.lstsq
    w[j] = N.linalg.pinv(X['tr'][:, :j]) @ Y['tr']
    R[j] = dict( (s, N.linalg.norm(X[s][:, :j] @ w[j]
                                   - Y[s])**2 / 2 / n)
                 for s in [ 'tr', 'te' ] )
```

Binomials, random walks, coin tosses, classification.

Binomials, random walks, coin tosses, classification.

Last lecture we saw *Hoeffding's Inequality*.

Theorem (Hoeffding's inequality). Suppose each draw from the distribution lies in the interval $[a, b]$. With probability at least $1 - \delta$ over an iid draw of $(z_i)_{i=1}^n$,

$$\mathbb{E}Z \leq \frac{1}{n} \sum_{i=1}^n z_i + (b - a) \sqrt{\frac{\ln(1/\delta)}{2n}}.$$

Binomials, random walks, coin tosses, classification.

Last lecture we saw *Hoeffding's Inequality*.

Theorem (Hoeffding's inequality). Suppose each draw from the distribution lies in the interval $[a, b]$. With probability at least $1 - \delta$ over an iid draw of $(z_i)_{i=1}^n$,

$$\mathbb{E}Z \leq \frac{1}{n} \sum_{i=1}^n z_i + (b - a) \sqrt{\frac{\ln(1/\delta)}{2n}}.$$


What does this mean?

Interpreting Hoeffding's inequality.

```
for i in range(k):  
    path = N.cumsum( N.random.randint(0, 2, n) * 2 - 1 )  
    plt.plot(path, color = 'black', alpha = 0.35)
```


Interpreting Hoeffding's inequality.

```
for i in range(k):  
    path = N.cumsum( N.random.randint(0, 2, n) * 2 - 1 )  
    plt.plot(path, color = 'black', alpha = 0.35)
```



Interpreting Hoeffding's inequality.


```
for i in range(k):  
    path = N.cumsum( N.random.randint(0, 2, n) * 2 - 1 )  
    plt.plot(path, color = 'black', alpha = 0.35)
```



Hoeffding says: for any **fixed** n , with probability $\geq 1 - 1/\sqrt{e}$,
$$\text{position} \leq \sqrt{n}.$$

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.

As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.

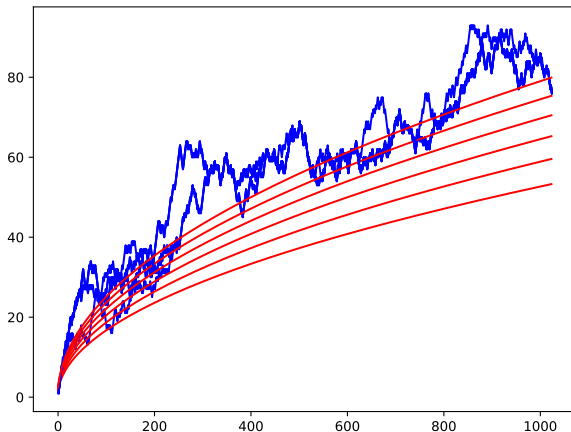
As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.

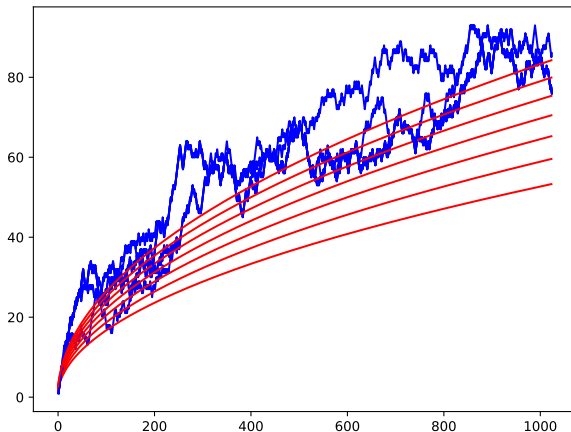
As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Hoeffding's inequality with more walks.

As paths are added, the upper bound **must** increase.



Question: how is the curve being rescaled?

Helpful tool: the union bound.

Given events (E_1, \dots, E_k) , the probability that *some* E_i occurs is

$$\Pr(E_1 \vee \dots \vee E_k) \leq \sum_i \Pr(E_i).$$

(Intuition: Venn diagram.)

Helpful tool: the union bound.

Given events (E_1, \dots, E_k) , the probability that *some* E_i occurs is

$$\Pr(E_1 \vee \dots \vee E_k) \leq \sum_i \Pr(E_i).$$

(Intuition: Venn diagram.)

A useful consequence.

Probability that *no* E_i occurs satisfies

$$\Pr(\neg E_1 \wedge \dots \wedge \neg E_k) \geq 1 - \sum_i \Pr(E_i).$$

Bounding *many* paths using Hoeffding.

Theorem (Hoeffding's inequality). Consider random variables (W_1, \dots, W_k) where $W_j := \frac{1}{n} \sum_{i=1}^n Z_{j,i}$ with $Z_{j,i} \in [a, b]$, and $Z_{j,i}$ are independent for fixed j , but may be *dependent* for fixed i . With probability at least $1 - \delta$,

$$\max_{j \in [k]} (\mathbb{E}(W_j) - W_j) \leq (b - a) \sqrt{\frac{\ln(k) + \ln(1/\delta)}{2n}}.$$

Bounding *many* paths using Hoeffding.

Theorem (Hoeffding's inequality). Consider random variables (W_1, \dots, W_k) where $W_j := \frac{1}{n} \sum_{i=1}^n Z_{j,i}$ with $Z_{j,i} \in [a, b]$, and $Z_{j,i}$ are independent for fixed j , but may be *dependent* for fixed i . With probability at least $1 - \delta$,

$$\max_{j \in [k]} (\mathbb{E}(W_j) - W_j) \leq (b - a) \sqrt{\frac{\ln(k) + \ln(1/\delta)}{2n}}.$$


► **Proof.** Define $\epsilon := \sqrt{\ln(k/\delta)/2n}$ and events

$$E_j := [\mathbb{E}(W_j) > W_j + \epsilon].$$

By Hoeffding, $\Pr(E_j) \leq \delta/k$, and by union bound

$$\Pr(\cap_j \neg E_j) \geq 1 - \sum_j \Pr(E_j) \geq 1 - \delta.$$


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.

Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#paths)}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#paths)}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#paths)}$.


Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.

Back to bounding random walks.



Question: how is the curve being rescaled?

Answer: $\sqrt{n \cdot \ln(\#\text{paths})}$.

Hoeffding and *one* classifier.

Let f be a *fixed* classifier. Define random variable

$$Z_i := \mathbb{1} [f(x_i) \neq y_i] .$$

Hoeffding gives: with probability at least $1 - \delta$,

$$\begin{aligned} \Pr[f(X) \neq Y] &= \mathbb{E}(Z_1) \\ &\leq \frac{1}{n} \sum_{i=1}^n Z_i + \sqrt{\frac{\ln(1/\delta)}{2n}} . \\ &\leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln(1/\delta)}{2n}} . \end{aligned}$$

Hoeffding and *many* classifiers.

This is **exactly** the “many paths” bound.

Hoeffding and *many* classifiers.

This is **exactly** the “many paths” bound.

Given classifiers \mathcal{F} , similarly: with probability $1 - \delta$, **every** $f \in \mathcal{F}$ satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

Complexity measures.

Given classifiers \mathcal{F} , with probability $1 - \delta$, **every** $f \in \mathcal{F}$ satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

Complexity measures.

Given classifiers \mathcal{F} , with probability $1 - \delta$, **every** $f \in \mathcal{F}$ satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

Generalization.

Given predictors \mathcal{F} , with probability at least $1 - \delta$, each $f \in \mathcal{F}$ satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{\mathcal{O}} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

Complexity measures.

Given classifiers \mathcal{F} , with probability $1 - \delta$, **every** $f \in \mathcal{F}$ satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

Generalization.

Given predictors \mathcal{F} , with probability at least $1 - \delta$, each $f \in \mathcal{F}$ satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{\mathcal{O}} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

Remark: holds for all of \mathcal{F} , in particular for *selected* $f \in \mathcal{F}$.

Complexity measures.

Given classifiers \mathcal{F} , with probability $1 - \delta$, **every** $f \in \mathcal{F}$ (including choice of an algorithm) satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

Complexity measures.

Given classifiers \mathcal{F} , with probability $1 - \delta$, **every** $f \in \mathcal{F}$ (including choice of an algorithm) satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$


where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

Example complexity measures.

- ▶ When $|\mathcal{F}| < \infty$, can use $\ln |\mathcal{F}|$.
- ▶ For classification, can use **VC dimension**.
- ▶ More generally, can use **Rademacher complexity**.

Binomials, random walks, coin tosses, classification.



Questions so far?

A few overfitting/generalization asides.

Aside: scientific experiments.

- ▶ **Standard scientific setup:** collect some data, try to fit various hypotheses to it.

Aside: scientific experiments.

- ▶ **Standard scientific setup:** collect some data, try to fit various hypotheses to it.
- ▶ **This is like checking multiple random walks!** The confidence intervals *must* grow with further hypotheses!

Aside: scientific experiments.

- ▶ **Standard scientific setup:** collect some data, try to fit various hypotheses to it.
- ▶ **This is like checking multiple random walks!** The confidence intervals *must* grow with further hypotheses!
- ▶ This observation is at the core of various “crises” in the application of statistics to science, and give the field of **adaptive data analysis**.

Aside: *covering number* complexities.

We have $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$;
why not **discretize** \mathcal{F} and apply?

Aside: *covering number* complexities.

We have $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$;
why not **discretize** \mathcal{F} and apply?

If $\|w - v\|$ small, then **predictions are similar**, meaning

$$|w^\top x - v^\top x| = |(w - v)^\top x| \leq \|w - v\| \|x\|.$$

Aside: *covering number* complexities.

We have $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$;
why not **discretize** \mathcal{F} and apply?

If $\|w - v\|$ small, then **predictions are similar**, meaning

$$|w^\top x - v^\top x| = |(w - v)^\top x| \leq \|w - v\| \|x\|.$$

Consequence: discretizing

$$\mathcal{F} := \{x \mapsto w^\top x : \|w\| \leq 1\}$$

is same as discretizing $\{w \in \mathbb{R}^d : \|w\| \leq 1\}$,

Aside: *covering number* complexities.

We have $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$;
why not **discretize** \mathcal{F} and apply?

If $\|w - v\|$ small, then **predictions are similar**, meaning

$$|w^\top x - v^\top x| = |(w - v)^\top x| \leq \|w - v\| \|x\|.$$

Consequence: discretizing

$$\mathcal{F} := \{x \mapsto w^\top x : \|w\| \leq 1\}$$

is same as discretizing $\{w \in \mathbb{R}^d : \|w\| \leq 1\}$,
which needs size $\mathcal{O}((1/\epsilon)^d)$, thus

$$\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(d \ln(1/\epsilon)).$$

Aside: VC dimension.

VC dimension gives a complexity measure for classifiers (binary output).

- ▶ **Definition:** the largest data set size (or ∞) which this function class (model) can label in all possible ways.

In earlier bounds, can plug in $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$.

Aside: VC dimension.

VC dimension gives a complexity measure for classifiers (binary output).

- ▶ **Definition:** the largest data set size (or ∞) which this function class (model) can label in all possible ways.

In earlier bounds, can plug in $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$.

Examples.

- ▶ **Linear separators:** d .
- ▶ **ReLU networks:** $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$.

Aside: neural networks.

- ▶ Our generalization bound scales with $\frac{\text{Complexity}(\mathcal{F})}{n}$.

Aside: neural networks.

- ▶ Our generalization bound scales with $\frac{\text{Complexity}(\mathcal{F})}{n}$.
- ▶ $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$, and ReLU networks have $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$.

Aside: neural networks.

- ▶ Our generalization bound scales with $\frac{\text{Complexity}(\mathcal{F})}{n}$.
- ▶ $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$, and ReLU networks have $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$.
- ▶ VC view says: ReLU networks need $\#\text{parameters} \cdot \#\text{layers} < n$.

Aside: neural networks.

- ▶ Our generalization bound scales with $\frac{\text{Complexity}(\mathcal{F})}{n}$.
- ▶ $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$, and ReLU networks have $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$.
- ▶ VC view says: ReLU networks need $\#\text{parameters} \cdot \#\text{layers} < n$.
- ▶ **False in practice** (for further discussion, google “deep learning rethinking generalization”).

Aside: neural networks.

- ▶ Our generalization bound scales with $\frac{\text{Complexity}(\mathcal{F})}{n}$.
- ▶ $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$, and ReLU networks have $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$.
- ▶ VC view says: ReLU networks need $\#\text{parameters} \cdot \#\text{layers} < n$.
- ▶ **False in practice** (for further discussion, google “deep learning rethinking generalization”).
- ▶ Can use other bounds: “small norm” property(?) of SGD and norm-based generalization.

Aside: Rademacher complexity.

VC is combinatorial/discrete;

Rademacher is a generalization that allows real-valued predictors.

- ▶ **Definition.** Let $(\epsilon_1, \dots, \epsilon_n)$ be **random sign** (“Rademacher”) random variables, meaning $\Pr[\epsilon_i = +1] = \Pr[\epsilon_i = -1] = 1/2$, and all are independent. Then

$$\text{Rad}(\mathcal{F}) := \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \epsilon_i f(x_i).$$

- ▶ **Intuition:** how well \mathcal{F} fits **random signs**.
- ▶ **Remark:** VC was **worst case** sign patterns.
- ▶ **Remark:** can plug in $\text{Complexity}(\mathcal{F}) \leq n \cdot \text{Rad}(\mathcal{F})^2$.

Aside: Rademacher complexity.

VC is combinatorial/discrete;

Rademacher is a generalization that allows real-valued predictors.

- ▶ **Definition.** Let $(\epsilon_1, \dots, \epsilon_n)$ be **random sign** (“Rademacher”) random variables, meaning $\Pr[\epsilon_i = +1] = \Pr[\epsilon_i = -1] = 1/2$, and all are independent. Then

$$\text{Rad}(\mathcal{F}) := \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \epsilon_i f(x_i).$$

- ▶ **Intuition:** how well \mathcal{F} fits **random signs**.
- ▶ **Remark:** VC was **worst case** sign patterns.
- ▶ **Remark:** can plug in $\text{Complexity}(\mathcal{F}) \leq n \cdot \text{Rad}(\mathcal{F})^2$.

Linear predictors with norm $\leq R$: then $\text{Rad}(\mathcal{F}) \leq R/\sqrt{n}$.

Neural networks with weight matrices (W_1, \dots, W_L) :

$$\text{Rad}(\mathcal{F}) = \tilde{O} \left((\text{gross stuff}) \cdot \prod_{i=1}^L \sigma_{\max}(W_i) \right).$$

Aside: ridge regression.

Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of λ in standard learning theory questions:

Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of λ in standard learning theory questions:

- **Representation:** not just a linear predictor, moreover $\mathcal{F}_\lambda := \{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}.$

Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of λ in standard learning theory questions:

- ▶ **Representation:** not just a linear predictor, moreover $\mathcal{F}_\lambda := \{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$.
- ▶ **Optimization:** as discussed in class, to achieve accuracy ϵ , gradient descent needs $\frac{\sigma_{\max}(X) + \lambda}{\sigma_{\min}(X) + \lambda} \ln(1/\epsilon)$ iterations.

Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of λ in standard learning theory questions:

- ▶ **Representation:** not just a linear predictor, moreover $\mathcal{F}_\lambda := \{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$.
- ▶ **Optimization:** as discussed in class, to achieve accuracy ϵ , gradient descent needs $\frac{\sigma_{\max}(X) + \lambda}{\sigma_{\min}(X) + \lambda} \ln(1/\epsilon)$ iterations.
- ▶ **Generalization:** via **Rademacher complexity** and above representation bound, get $\text{Complexity}(\mathcal{F}_\lambda) \leq 1/\lambda$.

Aside: online learning and the perceptron algorithm.

(Details in lecture.)

Summary (of overfitting/generalization).

Summary (of overfitting/generalization).

- ▶ **Intuition:** predictors/model too flexible \implies overfit (unless tons of data).
- ▶ **Rigorous form:** we have **generalization bounds**, namely bounds between training and test errors of the form

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right),$$

and we gave a few definitions of $\text{Complexity}(\mathcal{F})$.