

# Lecture 21 — Generative Adversarial Networks (GANs).

Alex Schwing and Matus Telgarsky

April 10, 2018

# Schedule for today.

1. Sampling with neural networks.
2. Original GAN formulation; Jensen-Shannon divergence.
3. Wasserstein GAN.
4. Some applications.

## Readings.

- ▶ (*Original GAN paper.*) Goodfellow, Ian J., Pouget-Abadie, Jean, Mirza, Mehdi, Xu, Bing, Warde-Farley, David, Ozair, Sherjil, Courville, Aaron C., and Bengio, Yoshua. “Generative adversarial nets”. *NIPS*, 2014.
- ▶ (Wasserstein GAN papers.)
  - ▶ Arjovsky, Martin, Chintala, Soumith, and Bottou, Leon. “Wasserstein generative adversarial networks”. *ICML*, 2017.
  - ▶ Gulrajani, Ishaan, Ahmed, Faruk, Arjovsky, Martin, Dumoulin, Vincent, and Courville, Aaron C. “Improved training of wasserstein gans”. *NIPS*, 2017.

Sampling with neural networks.

## Sampling with neural networks.

**Generative adversarial networks (GANs)** are a way to approximately sample from a distribution with neural networks.

1. Obtain a sample  $(x_i)_{i=1}^n$  from the distribution.
2. Use **adversarial training** to fit a neural network  $g$  to this sample. Specifically, perform the following alternating minimization.
  - 2.1 Generate fake sample  $(\tilde{x}_i)_{i=1}^m$  from  $g$ , and train another network  $f$ , the *adversary/discriminator*, to distinguish  $(x_i)_{i=1}^n$  and  $(\tilde{x}_i)_{i=1}^m$ .
  - 2.2 Now leave  $f$  fixed, and train  $g$  so that  $f$  no longer distinguishes fake and true samples.

## Sampling with neural networks.

**Generative adversarial networks (GANs)** are a way to approximately sample from a distribution with neural networks.

1. Obtain a sample  $(x_i)_{i=1}^n$  from the distribution.
2. Use **adversarial training** to fit a neural network  $g$  to this sample. Specifically, perform the following alternating minimization.
  - 2.1 Generate fake sample  $(\tilde{x}_i)_{i=1}^m$  from  $g$ , and train another network  $f$ , the *adversary/discriminator*, to distinguish  $(x_i)_{i=1}^n$  and  $(\tilde{x}_i)_{i=1}^m$ .
  - 2.2 Now leave  $f$  fixed, and train  $g$  so that  $f$  no longer distinguishes fake and true samples.

### Immediate questions.

1. What is the objective function?
2. How can networks generate data? (We'll discuss this now.)

# Generating data with neural networks.

Here is the approach used in GANs.

1. First sample  $z \sim \mu$  from some efficiently-sampled distribution  $\mu$ ; e.g., Gaussian or uniform.
2. Now output  $g(z)$ , where  $g$  is a neural network.

# Generating data with neural networks.

Here is the approach used in GANs.

1. First sample  $z \sim \mu$  from some efficiently-sampled distribution  $\mu$ ; e.g., Gaussian or uniform.
2. Now output  $g(z)$ , where  $g$  is a neural network.

**Notation:** new distribution is  $g\#\mu$ .

# Generating data with neural networks.

Here is the approach used in GANs.

1. First sample  $z \sim \mu$  from some efficiently-sampled distribution  $\mu$ ; e.g., Gaussian or uniform.
2. Now output  $g(z)$ , where  $g$  is a neural network.

**Notation:** new distribution is  $g\#\mu$ .


**Remark.**

- ▶ This is a way to sample, but how to estimate *probabilities*?
  - ▶ If we could invert  $g$  (in the set-valued sense), then  $\Pr[g^{-1}(S)]$  would make sense, where  $\Pr[\cdot]$  measures probability according to  $z$ .
  - ▶ Another approach is to use a kernel density estimate; this was used in the original GAN paper, but has a bad curse of dimension.
  - ▶ Basically, though: no one knows.




## Examples — mappings.

$g(x) = x$ , the identity function; target distribution  $\text{Uniform}([0, 1])$ .




## Examples — mappings.

$g(x) = x^2$ ; target density proportional to  $\frac{2}{\sqrt{x}}$  along  $(0, 1]$ .




## Examples — mappings.

$g(x)$  is inverse CDF of Gaussian; target distribution is Gaussian.




## Examples — univariate sampling.

Univariate sample, kernel density estimate (kde), GMM E-M.




## Examples — univariate sampling.

Univariate sample, kernel density estimate (kde), GAN kde.




## Examples — bivariate sampling.

Bivariate sample, GMM E-M.




## Examples — bivariate sampling.

Bivariate sample, kernel density estimate (kde).



## Examples — bivariate sampling.

Bivariate sample, GAN kde.



**Question:** how will this plot change with network capacity?



## Original GAN formulation.

Let's now discuss the choice of objective function.

## Original GAN formulation.

Let's now discuss the choice of objective function.

Recall the GAN setup:

- ▶ True sample  $(x_i)_{i=1}^n$ , fake sample  $(\tilde{x}_j)_{j=1}^m = (g(z_j))_{j=1}^m$ .
- ▶ We want these to look similar.

How to enforce similarity? One idea is to look at

$$\inf_{g \in \mathcal{G}} \sup_{f \in \mathcal{F}} \mathbb{E}(f(X)) - \mathbb{E}(f(g(Z)))$$

where  $\mathcal{F}$  is a big set of *adversaries/discriminators/critics*.

# Original GAN formulation.

Let's now discuss the choice of objective function.

Recall the GAN setup:

- ▶ True sample  $(x_i)_{i=1}^n$ , fake sample  $(\tilde{x}_j)_{j=1}^m = (g(z_j))_{j=1}^m$ .
- ▶ We want these to look similar.

How to enforce similarity? One idea is to look at

$$\inf_{g \in \mathcal{G}} \sup_{f \in \mathcal{F}} \mathbb{E}(f(X)) - \mathbb{E}(f(g(Z)))$$

where  $\mathcal{F}$  is a big set of *adversaries/discriminators/critics*.

## Remarks.

- ▶ Suppose  $\mathcal{F}$  is all polynomials. Then we get value 0 if true and fake data agree on all moments.
- ▶ If  $g$  is fixed, can optimize  $f$  over some neural nets  $\mathcal{F}$ . Similarly, can optimize  $g$  while holding  $f$  fixed.

## A transformation.

Now consider an adjustment, used in the original gan paper:

$$\inf_{g \in \mathcal{G}} \sup_{\substack{f \in \mathcal{F} \\ f: X \rightarrow (0,1)}} \frac{1}{n} \sum_{i=1}^n \ln(f(x_i)) + \frac{1}{m} \sum_{j=1}^m \ln(1 - f(g(z_j)))$$

### Remarks.

- Interpret  $f$  as a probability;  
e.g.,  $\Pr[x \text{ is fake}]$ .  
Then  $g$  is doing well if  $f = 1/2$  forced everywhere.  
(We'll do this systematically in a moment.)

# Original GAN formulation and algorithm.

Original GAN objective:

$$\inf_{g \in \mathcal{G}} \sup_{\substack{f \in \mathcal{F} \\ f: X \rightarrow (0,1)}} \frac{1}{n} \sum_{i=1}^n \ln (f(x_i)) + \frac{1}{m} \sum_{j=1}^m \ln (1 - f(g(z_j)))$$

Algorithm alternates these two steps:

1. Hold  $g$  fixed and optimize  $f$ . Specifically, generate a sample  $(\tilde{x}_j)_{j=1}^m = (g(z_j))_{j=1}^m$ , and approximately optimize

$$\sup_{\substack{f \in \mathcal{F} \\ f: X \rightarrow (0,1)}} \frac{1}{n} \sum_{i=1}^n \ln (f(x_i)) + \frac{1}{m} \sum_{j=1}^m \ln (1 - f(\tilde{x}_j)) .$$

2. Hold  $f$  fixed and optimize  $g$ . Specifically, generate  $(z_j)_{j=1}^m$  and approximately optimize

$$\inf_{g \in \mathcal{G}} \frac{1}{n} \sum_{i=1}^n \ln (f(x_i)) + \frac{1}{m} \sum_{j=1}^m \ln (1 - f(g(z_j))) .$$

## Some implementation issues.

Algorithm alternates these two steps:

1. Hold  $g$  fixed and optimize  $f$ . Specifically, generate a sample  $(\tilde{x}_j)_{j=1}^m = (g(z_j))_{j=1}^m$ , and approximately optimize

$$\sup_{\substack{f \in \mathcal{F} \\ f: X \rightarrow (0,1)}} \frac{1}{n} \sum_{i=1}^n \ln(f(x_i)) + \frac{1}{m} \sum_{j=1}^m \ln(1 - f(\tilde{x}_j)) .$$

2. Hold  $f$  fixed and optimize  $g$ . Specifically, generate  $(z_j)_{j=1}^m$  and approximately optimize

$$\inf_{g \in \mathcal{G}} \frac{1}{n} \sum_{i=1}^n \ln(f(x_i)) + \frac{1}{m} \sum_{j=1}^m \ln(1 - f(g(z_j))) .$$

### Remarks.

- ▶ Common practice: do many  $f$  ascents for each  $g$  descent.
- ▶ Training has all sorts of instabilities and heuristics fixes; e.g., **mode collapse** ( $g$  output a small set of training elements).

Original GAN objective function.

## Original GAN objective function.

We can interpret the original GAN as minimizing a divergence/distance between probability distributions.  
(In particular, with no adversary.)



# Original GAN objective function.

We can interpret the original GAN as minimizing a divergence/distance between probability distributions.  
(In particular, with no adversary.)

For this part, assume the following densities exist:

- ▶  $p_X$ , a density generating the data  $(x_i)_{i=1}^n$ .
- ▶  $p_Z$ , the density generating the input  $z$  to the generator network.
- ▶  $p_g$ , the density of the fake distribution of generator network  $g$ .

Some other technical caveats will be mentioned.

# Jensen-Shannon Divergence.

## Proposition.

- ▶ Given generator  $g$ , the optimal discriminator is  $f_g(x) := p_X(x)/(p_g(x)+p_X(x))$ .
- ▶ If discriminator set  $\mathcal{F}$  is all functions, the GAN objective is

$$\sup_{f \in \mathcal{F}} \mathbb{E}(\ln f(X)) + \mathbb{E}(\ln(1 - f(g(Z)))) = 2 \cdot \text{JS}(p_X, p_g) - \ln 4,$$

where JS is the *Jensen-Shannon divergence*:

defining  $p := (p_X + p_g)/2$ ,

$$\begin{aligned} 2 \cdot \text{JS}(p_X, p_g) &= \text{KL}(p_X, p) + \text{KL}(p_g, p) \\ &= \int p_X(x) \ln \frac{p_X(x)}{p(x)} dx + \int p_g(x) \ln \frac{p_g(x)}{p(x)} dx. \end{aligned}$$

- ▶ The optimal  $p_g$  satisfies  $p_g = p_X$ .

## Remarks.

- ▶ We can make this the starting point for the original GAN derivation.
  1. We want to find  $g$  so that  $\rho(p_g, p_X)$  is small for some reasonable notion of distance  $\rho$ .
  2. We choose Jensen-Shannon divergence for  $\rho$ . We could have chosen something else.
- ▶ From this perspective, the 2-player game view is a consequence which is used to derive a *training algorithm*.
- ▶ As with  $k$ -means/E-M, we massaged the objective to add another variable, and trained with alternating minimization.

## Proof of optimal discriminator.

Using the assumed densities  $p_X, p_Z, p_g$ ,

$$\begin{aligned} & \mathbb{E} \ln f(X) + \mathbb{E} \ln(1 - f(g(Z))) \\ &= \int \ln f(x) p_X(x) dx + \int \ln(1 - f(g(z))) p_Z(z) dz \\ &= \int \ln f(x) p_X(x) dx + \int \ln(1 - f(x)) p_g(x) dx. \\ &= \int \left( \ln f(x) p_X(x) + \ln(1 - f(x)) p_g(x) \right) dx. \end{aligned}$$

Since  $f$  can be any function, we can maximize it pointwise. Note  $r \mapsto a \ln(r) + b \ln(1 - r)$  is concave with maximum  $a/(a + b)$ .

Therefore, optimal discriminator satisfies  $f(x) = \frac{p_X(x)}{p_X(x) + p_g(x)}$ .

## Proof of alternate form.

Plugging this back in,

$$\begin{aligned} & \sup_{f \in \mathcal{F}} \mathbb{E} \ln f(X) + \mathbb{E} \ln(1 - f(g(Z))) \\ &= \sup_{f \in \mathcal{F}} \int \left( \ln f(x) p_X(x) + \ln(1 - f(x)) p_g(x) \right) dx. \\ &= \int \left( p_X(x) \ln \frac{p_X(x)}{p_X(x) + p_g(x)} + p_g(x) \ln \frac{p_g(x)}{p_X(x) + p_g(x)} \right) dx. \\ &= \int \left( p_X(x) \ln \frac{2p_X(x)}{p_X(x) + p_g(x)} + p_g(x) \ln \frac{2p_g(x)}{p_X(x) + p_g(x)} \right) dx - \ln 4 \\ &= \text{KL} \left( p_X, \frac{p_X + p_g}{2} \right) + \text{KL} \left( p_g, \frac{p_X + p_g}{2} \right) - \ln 4. \end{aligned}$$

## Technical remarks.

- ▶ This derivation is over the **true** distribution, not the sample! The sample induces a **discrete** distribution!
  - ▶ How to regularize/generalize?
  - ▶ The optimum of memorizing training set is trivial and doesn't need a GAN to train (just randomly sample the training set).
- ▶ The discriminator need only satisfy the stated equality with probability 1.
- ▶ By this derivation, mode collapse is not baked into the objective function; it is a side effect of training (e.g., non-convexity).
- ▶ The analysis needs  $\mathcal{F}$  to be all possible functions. But in general we use some restricted/regularized set of neural networks. What is the corresponding optimal discriminator?

## Technical remarks.

- ▶ The optimality condition on the generator,  $p_g = p_X$ , is a consequence of strict concavity of  $\ln$  inside Jensen-Shannon divergence.
- ▶ Similarly to the optimal discriminator equation, this only holds with probability 1, and requires the generator set  $\mathcal{G}$  to be everything (in general).
- ▶ There are many standard choices for the (restricted/regularized) set of generators. One is the “DCGAN”.

## Technical remarks.

- ▶ Rather than arguing pointwise, the optimal discriminator can be found by variants of “take gradient of objective, set to zero”.
  - ▶ A variant is needed because the optimization variable  $f$ , is an arbitrary function, not just a vector.
  - ▶ One such variant is mentioned in the homework. Namely, the “Euler-Lagrange equation”. This equation is for objectives that possess not only  $f$ , but its derivative. Namely, let  $\int L(x, f, f') dx$  denote the GAN objective. A corresponding optimality condition is

$$\frac{\partial L(x, f, f')}{\partial f} - \frac{d}{dx} \frac{\partial L(x, f, f')}{\partial f'} = 0.$$

Since  $f'$  does not appear, the second term is zero, and this equation becomes

$$\frac{p_x}{f} - \frac{p_g}{1-f} = 0,$$

which gives  $f = p_x / (p_g + p_x)$  as before.

- ▶ There are other versions of “take derivative and set to 0” in function spaces which can be used (e.g., ones without  $f'$ ).



Wasserstein GAN (WGAN).

## Wasserstein GAN (WGAN).

Let's build another GAN around another objective function.

Let's start from the 2-player perspective again.

## WGAN: 2-player formulation.

Recall the original abstract adversary optimization:

$$\inf_{g \in \mathcal{G}} \sup_{f \in \mathcal{F}} \mathbb{E}(f(X)) - \mathbb{E}(f(g(Z))).$$

- ▶ The Wasserstein GAN imposes a specific constraint:  
 $\|f\|_{\text{Lip}} \leq 1$ : the set of functions with Lipschitz constant less than 1 (this means  $\sup_{x \neq y} \frac{f(x) - f(y)}{\|x - y\|} \leq 1$ , and for differentiable functions means  $|f'| \leq 1$ ).
- ▶ In practice,  $\mathcal{F}$  is a set of neural nets plus some regularization/constraints aiming to enforce  $\|f\|_{\text{Lip}} \leq 1$ . It's not clear how best to do this; the two WGAN papers mentioned in the readings do it differently.

## WGAN: non-adversarial objective.

Due to the Lipschitz constraint, the objective is the Wasserstein-1 distance; here is a special case for densities  $p_g$ ,  $p_X$  as before:

$$\begin{aligned} W(p_g, p_X) &= \sup_{\|f\|_{\text{Lip}} \leq 1} \left( \int f(x) p_X(x) dx - \int f(x) p_g(x) dx \right) \\ &= \inf \left\{ \int \|x - h(x)\| p_g(x) dx : h \# p_g = p_X \right\} \end{aligned}$$

- ▶ This is called the **Kantorovich-Rubinstein** duality.
- ▶ The meaning of the last expression is: if we sample  $x \sim p_g$  and output  $h(x)$ , this is the same as sampling  $p_X$ .
- ▶ This distance is also called “Earth mover’s distance” because it can be interpreted as moving the mass from one distribution into the shape of another, so that particles are moved a minimal distance. (Pictures in class.)

# Summary.

- ▶ Generating random samples with neural networks.
- ▶ Original GAN.
  - ▶ Objective function.
  - ▶ Alternating optimization (min/max).
  - ▶ The optimal choices (supposing all possible discriminators/generators).
- ▶ Wasserstein GAN
  - ▶ Objective function.