

Machine Learning

A. G. Schwing & M. Telgarsky

University of Illinois at Urbana-Champaign, 2018

L22: Autoregressive Methods (RNNs/LSTMs/GRUs)

Goals of this lecture

Goals of this lecture

- Getting to know Recurrent Neural Nets (RNNs)

Goals of this lecture

- Getting to know Recurrent Neural Nets (RNNs)
- Getting to know Long short term memory (LSTM)

Goals of this lecture

- Getting to know Recurrent Neural Nets (RNNs)
- Getting to know Long short term memory (LSTM)
- Getting to know Gated recurrent unit (GRU)

Goals of this lecture

- Getting to know Recurrent Neural Nets (RNNs)
- Getting to know Long short term memory (LSTM)
- Getting to know Gated recurrent unit (GRU)
- Seeing how to apply them

Goals of this lecture

- Getting to know Recurrent Neural Nets (RNNs)
- Getting to know Long short term memory (LSTM)
- Getting to know Gated recurrent unit (GRU)
- Seeing how to apply them

Reading Material

- Goodfellow et al.; Deep Learning; Chapter 10
- Papers cited on the slides

Pixel Recurrent Neural Networks



Recap: Our models so far

Recap: Our models so far

- Discriminative

Recap: Our models so far

- Discriminative

$$p(\mathbf{y}|x)$$

Recap: Our models so far

- Discriminative

$$p(\mathbf{y}|x)$$

- Generative

Recap: Our models so far

- Discriminative

$$p(\mathbf{y}|x)$$

- Generative

$$p(x)$$


Recap: Our models so far

- Discriminative

$$p(\mathbf{y}|x)$$

- Generative

$$p(x)$$




Recap: Our models so far

- Discriminative

$$p(\mathbf{y}|x)$$

- Generative

$$p(x)$$



What's missing?

More flexibility regarding inputs and outputs:

More flexibility regarding inputs and outputs:

- Sequences of inputs

More flexibility regarding inputs and outputs:

- Sequences of inputs
- Sequences of outputs

More flexibility regarding inputs and outputs:

- Sequences of inputs
- Sequences of outputs

Length of sequences may vary

More flexibility regarding inputs and outputs:

- Sequences of inputs
- Sequences of outputs

Length of sequences may vary

one to one




More flexibility regarding inputs and outputs:

- Sequences of inputs
- Sequences of outputs

Length of sequences may vary

one to one one to many




More flexibility regarding inputs and outputs:

- Sequences of inputs
- Sequences of outputs

Length of sequences may vary

one to one one to many many to one




More flexibility regarding inputs and outputs:

- Sequences of inputs
- Sequences of outputs

Length of sequences may vary

one to one one to many many to one many to many




More flexibility regarding inputs and outputs:


- Sequences of inputs
- Sequences of outputs

Length of sequences may vary

one to one one to many many to one many to many many to many




Recurrent Neural Nets (RNNs)



Recurrent Neural Nets (RNNs)


- input depends on previous output



Recurrent Neural Nets (RNNs)

- input depends on previous output


$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$



Recurrent Neural Nets (RNNs)

- input depends on previous output

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$




Recurrent Neural Nets (RNNs)

- input depends on previous output

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$


Applications:



Recurrent Neural Nets (RNNs)

- input depends on previous output

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$




Applications:

- Natural language processing

Recurrent Neural Nets (RNNs)

- input depends on previous output

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$




Applications:

- Natural language processing
- Speech recognition

Recurrent Neural Nets (RNNs)

- input depends on previous output

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$




Applications:

- Natural language processing
- Speech recognition
- Image processing

Recurrent Neural Nets (RNNs)

- input depends on previous output

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$




Applications:


- Natural language processing
- Speech recognition
- Image processing
- Video processing

Important concept: Parameter sharing

Important concept: Parameter sharing




Important concept: Parameter sharing




$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Important concept: Parameter sharing





$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Important concept: Parameter sharing



$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\ y^{(t)} &= g(h^{(t)})\end{aligned}$$


Important concept: Parameter sharing




$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\ y^{(t)} &= g(h^{(t)})\end{aligned}$$

unfolded/unrolled network
performs identical operations
easier to understand

General structure for recurrence:




General structure for recurrence:



Mathematical description in general:


General structure for recurrence:



Mathematical description in general:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

General structure for recurrence:



Mathematical description in general:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\ y^{(t)} &= g(h^{(t)})\end{aligned}$$

Note that f and g are independent of time

What are f and g ?

What are f and g ?

Any differentiable function can be used

What are f and g ?

Any differentiable function can be used

Useful functions:

What are f and g ?

Any differentiable function can be used

Useful functions:

- Original recurrent nets

What are f and g ?

Any differentiable function can be used

Useful functions:

- Original recurrent nets
- LSTM nets

What are f and g ?

Any differentiable function can be used

Useful functions:

- Original recurrent nets
- LSTM nets
- GRU nets

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

Specifically:

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

Specifically:

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

Specifically:

$$h^{(t)} = \sigma_h(W_{hx}x^{(t)} + W_{hh}h^{(t-1)} + w_{hb})$$

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

$$y^{(t)} = g(h^{(t)})$$

Specifically:

$$h^{(t)} = \sigma_h(W_{hx}x^{(t)} + W_{hh}h^{(t-1)} + w_{hb})$$

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

$$y^{(t)} = g(h^{(t)})$$

Specifically:

$$h^{(t)} = \sigma_h(W_{hx}x^{(t)} + W_{hh}h^{(t-1)} + w_{hb})$$

$$y^{(t)} = \sigma_y(W_{yh}h^{(t)} + w_{yb})$$

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

$$y^{(t)} = g(h^{(t)})$$

Specifically:

$$h^{(t)} = \sigma_h(W_{hx}x^{(t)} + W_{hh}h^{(t-1)} + w_{hb})$$

$$y^{(t)} = \sigma_y(W_{yh}h^{(t)} + w_{yb})$$

What is σ_h and σ_y ?

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

$$y^{(t)} = g(h^{(t)})$$

Specifically:

$$h^{(t)} = \sigma_h(W_{hx}x^{(t)} + W_{hh}h^{(t-1)} + w_{hb})$$

$$y^{(t)} = \sigma_y(W_{yh}h^{(t)} + w_{yb})$$

What is σ_h and σ_y ?

Activation functions:

Original recurrent nets (Elman network):

(Jordan network is slightly different)

Generally:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}, \mathbf{w})$$

$$y^{(t)} = g(h^{(t)})$$

Specifically:

$$h^{(t)} = \sigma_h(W_{hx}x^{(t)} + W_{hh}h^{(t-1)} + w_{hb})$$

$$y^{(t)} = \sigma_y(W_{yh}h^{(t)} + w_{yb})$$

What is σ_h and σ_y ?

Activation functions: tanh, sigmoid

Affine transformations and point-wise non-linearity

What are the problems?

Problems with classical recurrent neural nets:

Problems with classical recurrent neural nets:


- Vanishing gradients

Problems with classical recurrent neural nets:

- Vanishing gradients
- Long term dependency

Problems with classical recurrent neural nets:

- Vanishing gradients
- Long term dependency



Long short term memory (LSTM)

Long short term memory (LSTM)

- Particular functional relation

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Specifically: (\circ denotes Hadamard product; σ is activation function)

$$i^{(t)} = \sigma_i(\mathbf{W}_{ix}x^{(t)} + \mathbf{W}_{ih}h^{(t-1)} + \mathbf{w}_{bi}) \quad \text{Input gate}$$

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Specifically: (\circ denotes Hadamard product; σ is activation function)

$$\begin{aligned}i^{(t)} &= \sigma_i(\mathbf{W}_{ix}x^{(t)} + \mathbf{W}_{ih}h^{(t-1)} + \mathbf{w}_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(\mathbf{W}_{fx}x^{(t)} + \mathbf{W}_{fh}h^{(t-1)} + \mathbf{w}_{bf}) && \text{Forget gate}\end{aligned}$$

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned} h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\ y^{(t)} &= g(h^{(t)}) \end{aligned}$$

Specifically: (\circ denotes Hadamard product; σ is activation function)

$$\begin{aligned} i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\ f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\ o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \end{aligned}$$

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Specifically: (\circ denotes Hadamard product; σ is activation function)

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell}\end{aligned}$$

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Specifically: (\circ denotes Hadamard product; σ is activation function)

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell}\end{aligned}$$

Long short term memory (LSTM)

- Particular functional relation
- Shown to better capture long-term dependencies
- Shown to address the vanishing gradient problem

Generally:

$$\begin{aligned}h^{(t)} &= f(h^{(t-1)}, x^{(t)}, \mathbf{w}) \\y^{(t)} &= g(h^{(t)})\end{aligned}$$

Specifically: (\circ denotes Hadamard product; σ is activation function)

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

• $i^{(t)}$:

Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:


- $i^{(t)}$: Does $x^{(t)}$ matter?

Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

- $i^{(t)}$: Does $x^{(t)}$ matter?




Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

- $i^{(t)}$: Does $x^{(t)}$ matter?
- $f^{(t)}$:




Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

- $i^{(t)}$: Does $x^{(t)}$ matter?
- $f^{(t)}$: Should $c^{(t-1)}$ be forgotten?




Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

- $i^{(t)}$: Does $x^{(t)}$ matter?
- $f^{(t)}$: Should $c^{(t-1)}$ be forgotten?
- $o^{(t)}$:




Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

- $i^{(t)}$: Does $x^{(t)}$ matter?
- $f^{(t)}$: Should $c^{(t-1)}$ be forgotten?
- $o^{(t)}$: How much $c^{(t)}$ should be exposed?




Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:

- $i^{(t)}$: Does $x^{(t)}$ matter?
- $f^{(t)}$: Should $c^{(t-1)}$ be forgotten?
- $o^{(t)}$: How much $c^{(t)}$ should be exposed?
- $\tilde{c}^{(t)}$:




Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:


- $i^{(t)}$: Does $x^{(t)}$ matter?
- $f^{(t)}$: Should $c^{(t-1)}$ be forgotten?
- $o^{(t)}$: How much $c^{(t)}$ should be exposed?
- $\tilde{c}^{(t)}$: Compute new memory



Equations:

$$\begin{aligned}i^{(t)} &= \sigma_i(W_{ix}x^{(t)} + W_{ih}h^{(t-1)} + w_{bi}) && \text{Input gate} \\f^{(t)} &= \sigma_f(W_{fx}x^{(t)} + W_{fh}h^{(t-1)} + w_{bf}) && \text{Forget gate} \\o^{(t)} &= \sigma_o(W_{ox}x^{(t)} + W_{oh}h^{(t-1)} + w_{bo}) && \text{Output/Exposure gate} \\\tilde{c}^{(t)} &= \sigma_c(W_{cx}x^{(t)} + W_{ch}h^{(t-1)} + w_{bc}) && \text{New memory cell} \\c^{(t)} &= f^{(t)} \circ c^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)} && \text{Final memory cell} \\h^{(t)} &= o^{(t)} \circ \sigma_h(c^{(t)})\end{aligned}$$

Intuition:




Long short term memory (LSTM):

Long short term memory (LSTM):

- Can be interpreted as a block in a neural net

Long short term memory (LSTM):

- Can be interpreted as a block in a neural net



Gated recurrent unit (GRU):

Gated recurrent unit (GRU):

- Performance similar to LSTM

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Equations: (\circ denotes Hadamard product)

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Equations: (\circ denotes Hadamard product)

$$z^{(t)} = \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) \quad \text{Update gate}$$

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Equations: (\circ denotes Hadamard product)

$$z^{(t)} = \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz})$$

Update gate

$$r^{(t)} = \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br})$$

Reset gate

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Equations: (\circ denotes Hadamard product)

$$\begin{aligned}
 z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\
 r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\
 \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell}
 \end{aligned}$$

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Equations: (\circ denotes Hadamard product)

$$\begin{aligned}
 z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\
 r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\
 \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\
 h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state}
 \end{aligned}$$

Gated recurrent unit (GRU):

- Performance similar to LSTM
- Fewer parameters compared to LSTM (no output gate)

Equations: (\circ denotes Hadamard product)

$$\begin{aligned}
 z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\
 r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\
 \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\
 h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state}
 \end{aligned}$$

Can again be interpreted as a block in the computation graph

Equations: (\circ denotes Hadamard product)

$$\begin{aligned} z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\ r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\ \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\ h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state} \end{aligned}$$

Intuition:

Equations: (\circ denotes Hadamard product)

$$\begin{aligned} z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\ r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\ \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\ h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state} \end{aligned}$$

Intuition:

• $r^{(t)}$:

Equations: (\circ denotes Hadamard product)

$$\begin{aligned} z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\ r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\ \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\ h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state} \end{aligned}$$

Intuition:

- $r^{(t)}$: Include $h^{(t-1)}$ in new memory?

Equations: (\circ denotes Hadamard product)

$$\begin{aligned} z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\ r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\ \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\ h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state} \end{aligned}$$

Intuition:

- $r^{(t)}$: Include $h^{(t-1)}$ in new memory?
- $z^{(t)}$:

Equations: (\circ denotes Hadamard product)


$$\begin{aligned} z^{(t)} &= \sigma_z(W_{zx}x^{(t)} + W_{zh}h^{(t-1)} + w_{bz}) && \text{Update gate} \\ r^{(t)} &= \sigma_r(W_{rx}x^{(t)} + W_{rh}h^{(t-1)} + w_{br}) && \text{Reset gate} \\ \tilde{h}^{(t)} &= \sigma_h(W_{hx}x^{(t)} + W_{rwh}(r^{(t)} \circ h^{(t-1)}) + w_{bh}) && \text{New memory cell} \\ h^{(t)} &= (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)} && \text{Hidden state} \end{aligned}$$

Intuition:


- $r^{(t)}$: Include $h^{(t-1)}$ in new memory?
- $z^{(t)}$: How much $h^{(t-1)}$ in next state?

Recurrent nets generally:

Recurrent nets generally:




Recurrent nets generally:



Other variants:


Recurrent nets generally:



Other variants:

- Bi-directional LSTMs [Schuster&Paliwal (1997), Graves&Schmidhuber (2005)]

Recurrent nets generally:



Other variants:

- Bi-directional LSTMs [Schuster&Paliwal (1997), Graves&Schmidhuber (2005)]
- Continuous time RNNs

How do we learn the parameters in the network?

How do we learn the parameters in the network?

$$p(y_1, \dots, y_T) = \prod_{i=1}^T p(y_i | y_1, \dots, y_{i-1})$$

How do we learn the parameters in the network?

$$p(y_1, \dots, y_T) = \prod_{i=1}^T p(y_i | y_1, \dots, y_{i-1})$$

Maximum likelihood specifies loss function

How do we learn the parameters in the network?

$$p(y_1, \dots, y_T) = \prod_{i=1}^T p(y_i | y_1, \dots, y_{i-1})$$

Maximum likelihood specifies loss function

Relation to structured models?


Training via gradient descent:

Training via gradient descent:

- How?
- What order?
- What information do we need to store?


Training via gradient descent:

- How?
- What order?
- What information do we need to store?



Training via gradient descent:

- How?
- What order?
- What information do we need to store?




Backpropagation through time (BPTT)


Pixel Recurrent Neural Networks



PixelRNN model (Autoregressive model):



Context



Multi-scale context

Generative models overview:

Generative models overview:

- Variational Auto-encoders (VAEs):

Generative models overview:

- Variational Auto-encoders (VAEs):
- Generative Adversarial Nets (GANs):

Generative models overview:

- Variational Auto-encoders (VAEs):
- Generative Adversarial Nets (GANs):
- Autoregressive models (RNNs):

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro:
 - ▶ Con:
- Generative Adversarial Nets (GANs):
 - ▶ Pro:
 - ▶ Con:
- Autoregressive models (RNNs):
 - ▶ Pro:
 - ▶ Con:

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con:
- Generative Adversarial Nets (GANs):
 - ▶ Pro:
 - ▶ Con:
- Autoregressive models (RNNs):
 - ▶ Pro:
 - ▶ Con:

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con: slightly blurry examples
- Generative Adversarial Nets (GANs):
 - ▶ Pro:
 - ▶ Con:
- Autoregressive models (RNNs):
 - ▶ Pro:
 - ▶ Con:

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con: slightly blurry examples
- Generative Adversarial Nets (GANs):
 - ▶ Pro: generate sharp images
 - ▶ Con:
- Autoregressive models (RNNs):
 - ▶ Pro:
 - ▶ Con:

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con: slightly blurry examples
- Generative Adversarial Nets (GANs):
 - ▶ Pro: generate sharp images
 - ▶ Con: difficult to optimize (unstable) [at least early variants]
- Autoregressive models (RNNs):
 - ▶ Pro:
 - ▶ Con:

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con: slightly blurry examples
- Generative Adversarial Nets (GANs):
 - ▶ Pro: generate sharp images
 - ▶ Con: difficult to optimize (unstable) [at least early variants]
- Autoregressive models (RNNs):
 - ▶ Pro: stable training & good likelihoods
 - ▶ Con:

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con: slightly blurry examples
- Generative Adversarial Nets (GANs):
 - ▶ Pro: generate sharp images
 - ▶ Con: difficult to optimize (unstable) [at least early variants]
- Autoregressive models (RNNs):
 - ▶ Pro: stable training & good likelihoods
 - ▶ Con: inefficient sampling & no low-dimensional codes

Generative models overview:

- Variational Auto-encoders (VAEs):
 - ▶ Pro: probabilistic graphical model interpretation
 - ▶ Con: slightly blurry examples
- Generative Adversarial Nets (GANs):
 - ▶ Pro: generate sharp images
 - ▶ Con: difficult to optimize (unstable) [at least early variants]
- Autoregressive models (RNNs):
 - ▶ Pro: stable training & good likelihoods
 - ▶ Con: inefficient sampling & no low-dimensional codes

Very active research area

Quiz:

Quiz:

- Describe the prediction process for an RNN?

Quiz:

- Describe the prediction process for an RNN?
- Describe the training process for RNNs?

Quiz:

- Describe the prediction process for an RNN?
- Describe the training process for RNNs?
- Contrast generative modeling techniques?

Important topics of this lecture

Important topics of this lecture

- Getting to know RNNs and its variants

Important topics of this lecture

- Getting to know RNNs and its variants
- Getting to know their use

Important topics of this lecture

- Getting to know RNNs and its variants
- Getting to know their use
- Contrasting RNNs to generative models

Important topics of this lecture

- Getting to know RNNs and its variants
- Getting to know their use
- Contrasting RNNs to generative models

Next up:

Reinforcement learning